

# Data Management

Leonardo Maria Carrozzo

2022/2023



# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	The relational data model . . . . .	1
1.2	Relational algebra . . . . .	3
1.2.1	Union . . . . .	3
1.2.2	Difference . . . . .	3
1.2.3	Cartesian Product . . . . .	4
1.2.4	Projection Operation . . . . .	4
1.2.5	Selection Operation . . . . .	4
1.2.6	Relational Algebra Expression . . . . .	5




# 1 Introduction

The course is based on the following topics:

- **The structure of a Data Base Management System (DBMS):** Relational data and queries, Buffer manager;
- **Transaction management:** The concept of transaction, Concurrency management;
- **Crash management:** Classification of failures, Recovery;
- **Data Warehousing:** Data warehousing architectures and operators, Data warehousing design;
- **NoSQL databases:** Document-based databases (such as MongoDB), Graph databases OLAP vs OLTP (such as Neo4j);
- **Physical structures for data bases:** File organizations for data base management, Principles of physical database design;
- **Query processing:** Evaluation of relational algebra operators, Fundamentals of query optimization;

## 1.1 The relational data model

A database in the Relational Model is a **set of tables** (or **relations**). Each **table** is a **set of rows** (or **tuples**). Each one with the **same set of columns** (or **attributes**).



$v$  is the value of the corresponding column and row. The attributes B and C form a **superkey**.

We have then:

- **Integrity constraint:** a rule at the level of the schema that all the rows must respect;
- **Superkey:** there cannot exist two or more rows that have the same value as the combination of multiple attributes;

- **Key:** attribute in a table;
- **Foreign key:** attributes in a table are a reference of another table;
- **Primary key:** special key that doesn't allow null values (a null value is a special value that says that the value is missing).

<i>A</i>	<i>B</i>	<i>C</i>
		$c_1$
		$c_2$
		$c_3$
		$c_3$

$T_1$  Unordered set

<i>D</i>	<i>E</i>	<i>F</i>
$c_0$		
$c_1$		
$c_2$		
$c_3$		

$T_2$  Must not miss any key

<i>A</i>	<i>B</i>	<i>C</i>
$a_1$	$b_1$	$c_1$
$a_1$	$b_2$	$c_2$
$a_2$	$b_2$	$c_3$
<i>null</i>	$b_2$	$c_3$

$T_1$  Unordered set

<i>D</i>	<i>E</i>	<i>F</i>
$c_0$		
$c_1$		
$c_2$		
$c_3$		

$T_2$  Must not miss any key

We have a "no predicate" on the null value: never equal and never different, comparison is always false.

As we have said, the Relational Data Model uses the mathematical concept of a relation as the formalism for describing and representing data. A relation is a subset of a cartesian product of sets. A relation can be considered as a "table" with rows and columns.

Codd introduced two different query languages for the relational data model:

- **Relational Algebra**, which is a procedural language. It is an algebraic formalism in which queries are expressed by applying a sequence of operations to relations.
- **Relational Calculus**, which is a declarative language. It is a logical formalism in which queries are expressed as formulas of first-order logic.

**Codd's Theorem:** Relational Algebra and Relational Calculus are essentially equivalent in terms of expressive power.

DBMSs are based on **SQL**, a hybrid of a procedural and a declarative language that combines features from both relational algebra and relational calculus.

## 1.2 Relational algebra

The operators of Relational Algebra can be divided into two groups:

- Three standard set-theoretic binary operations:
  - Union
  - Difference
  - Cartesian Product
- Two special unary operations on relations:
  - Projection
  - Selection

The Relational Algebra consists of all expressions obtained by combining these five basic operations in syntactically correct ways.

- In relational algebra, both arguments of the union and the difference must be relations of the same arity.
- In SQL, there is the additional requirement that the corresponding attributes must have the same data type.
- However, the corresponding attributes need not have the same names; the corresponding attribute in the result can be renamed arbitrarily.

### 1.2.1 Union

- Takes in input two k-ary relations R and S, for some k.
- Gives in output the k-ary relation  $R \cup S$ , where:

$$R \cup S = \{(a_1, \dots, a_k) : (a_1, \dots, a_k) \text{ is in } R \text{ or } (a_1, \dots, a_k) \text{ is in } S\}$$

### 1.2.2 Difference

- Takes in input two k-ary relations R and S, for some k.
- Gives in output the k-ary relation  $R - S$ , where:

$$R - S = \{(a_1, \dots, a_k) : (a_1, \dots, a_k) \text{ is in } R \text{ and } (a_1, \dots, a_k) \text{ is not in } S\}$$

### 1.2.3 Cartesian Product

- Take in input an m-ary relation R and an n-ary relation S.
- Gives in output the (m+n)-ary relation  $R \times S$ , where:

$$R \times S = \{(a_1, \dots, a_m, b_1, \dots, b_n) : (a_1, \dots, a_m) \text{ is in } R \text{ and } (b_1, \dots, b_n) \text{ is in } S\}$$

Note:  $|R \times S| = |R| \times |S|$

Let's see an example:

Emp	Dept
Rossi	A
Neri	B
Bianchi	B

Employee

Dept	Char
A	Mori
B	Bruni

Dept

Emp	Dept	Code	Chair
Rossi	A	A	Mori
Rossi	A	B	Bruni
Neri	B	A	Mori
Neri	B	B	Bruni
Bianchi	B	A	Mori
Bianchi	B	B	Bruni

Employee  $\times$  Dept

### 1.2.4 Projection Operation

Given a table R, we want to rearrange the order of the columns and/or suppress/rename some columns.

Projection is a family of unary operations of the form:

$$\pi_{\langle \text{attribute list} \rangle}(\langle \text{relation name} \rangle) \text{ or } \text{PROJ}_{\langle \text{attribute list} \rangle}(\langle \text{relation name} \rangle)$$

When the projection is applied to a relation R, it removes all columns whose attributes do not appear in the  $\langle \text{attribute list} \rangle$  (we assume that an attribute can appear only once in the list).

The remaining columns may be re-arranged (and also renamed by means of the notation  $a \leftarrow b$ ) according to the order and name in the  $\langle \text{attribute list} \rangle$ .

Any duplicate rows are eliminated.

### 1.2.5 Selection Operation

Selection is a family of unary operations of the form:

$$\sigma_{\theta}(R) \text{ or } \text{SEL}_{\theta}(R)$$



where  $R$  is a relation and  $\theta$  is a condition that can be applied as a test to each row of  $R$ . When a selection operation is applied to  $R$ , it returns the subset of  $R$  consisting of all rows that satisfy the condition  $\theta$ .

Here are some examples:  $\sigma_{A=10}(T)$  or  $\sigma_{(A=10 \text{ or } B>20) \text{ and } C \text{ is not null}}(T)$ . Where  $A = 10$  is a boolean expression.  $T$  might be an expression or a table.

We have two special predicates:

- is null
- is not null

A condition in the selection operation is an expression built up from:

- Comparison operators  $=, <, >, \neq, \leq, \geq$  applied to operands that are constants or attribute names or component numbers. (These are the basic (atomic) clauses of the conditions).
- The boolean logic operators  $\wedge, \vee, :$  applied to basic clauses.

### 1.2.6 Relational Algebra Expression

A relational algebra expression is an expression obtained from relation schemas using union, difference, cartesian product, projection, and selection.